

Fig. 7 | Deployment interpretation framework

Audit outcomes are mapped to validation, recalibration, retraining or no-deployment actions.

Raw high + background-centered high + acceptable calibration	candidate_for_controlled_deployment	Proceed only with local calibration check, locked thresholds, and ongoing drift monitoring
Raw high + background-centered high + poor calibration	ranking_only_recalibrate_before_clinical_use	Use only as a ranking signal until calibration, thresholds, and clinical operating point
Raw high + background-centered low	background_dependent_retrain_locally	Do not deploy as portable focal-drug model; validate locally and retrain or recalibrate
Raw high + underpowered matching	insufficient_matched_evidence	Collect more matched support or additional external labels before making a deployment claim
Raw low + background-centered low	not_deployment_ready	Do not deploy; use for research triage or redesign the model/evaluation target.